

Skills-Graph Architecture

A unified reference for importing Bernd's experience records into a **Hyper-graph-of-Thought** in Neo4j and exposing it through an **MCP** service that local or cloud-hosted LLM agents can query.

1 Scope & Goals

- **Author** Bernd Prager
- **Revision** v1.0 · 2025-05-15
- **Purpose** Define components, data flows, schemas, and operational guidelines so that:
 - Markdown records (jobs, extras, certifications) → **hyper-graph** (Neo4j)
 - Graph → **MCP server API**
 - LLM/agent stack (Ollama + LangChain) can answer skill-centric queries.

2 Logical View

3 Data Sources

Folder	Type	Example File	Primary Entities
docs/jobs/	Job experience	EPAM.md	Role, Project, Skill, Tool
docs/extras/	Extra-professional	Ext_WJD.md	Activity, Skill
docs/certs/	Certifications	certifications.md	Certification, Skill

Registry Table (doc_registry)

Column	Type	Description
doc_id	TEXT PK	stem of filename
hash	CHAR(64)	SHA-256 checksum
last_ingested	DATETIME	UTC timestamp

4 Ingestion Worker

- **Language** Python 3.11
- **Key libs** langchain-community, neo4j-driver, faiss-cpu, pyyaml, python-multipart.

4.1 Steps per document

1. **Hash check** Skip if unchanged.
2. **Chunk** ~1500 tokens with overlap = 200.
3. **Embeddings** nomic-embed-text → FAISS index (shared).
4. **LLM IE** gemma3:12b prompt with known skills/tools.
5. **Dedup** similarity lookup (≥ 0.83 FTS OR ≥ 0.88 embed).
6. **Cypher MERGE** nodes + rels.
7. **Hyperedge build** hash(sorted node-ids) → create/update.

4.2 Config File

Store schema & prompt hints in **graph_schema.yaml** (see separate file).

5 Graph Schema (Neo4j)

Refer to **graph_schema.yaml** for machine-readable detail.

- **Core labels** Person, Role, Organization, Project, Activity, Certification, Skill, Tool, Topic, Hyperedge.
 - **Key rels** HAS_ROLE, WORKED_AT, CONTRIBUTED_TO, USED_IN, SHOWCASED_IN, COVERS_TOPIC, OWNS_CERT, SUPPORTS_SKILL, CONNECTS.
 - **Indexes**
 - CREATE CONSTRAINT person_name IF NOT EXISTS ON (p:Person) ASSERT p.name IS UNIQUE;
 - CREATE FULLTEXT INDEX skill_name IF NOT EXISTS FOR (s:Skill) ON EACH [s.name];
-

6 Infrastructure Topology

Host	Stack	Ports
odin (Ubuntu 22 LTS, Ryzen 9)	Neo4j 5.15 (Docker)	7474/7687
idem	Ollama 0.1.x (models in /var/lib/ollama)	11434

Host	Stack	Ports
idem	Ingestion Worker (systemd unit)	–
idem	FastAPI MCP server	8000

Note RTX 2060 (6 GB VRAM) runs `gemma3:12b Q4_0`; bigger models spill to RAM.

7 MCP Server

- **Framework** FastAPI + LangChain Graph
- **Auth** Bearer JWT (future: OIDC)
- **Endpoints**
 - POST `/query` → JSON {prompt, agent=“graph-rag”}
 - POST `/skill_matrix` → returns CSV of skills vs. evidence nodes
- **Agent Types**
 - **GraphRAGAgent** → augments prompt with Cypher results
 - **SkillLensAgent** → tailored skill queries (`MATCH (s:Skill)<-[:SHOWCASED_IN]-(n)...`)

8 Observability

- **Neo4j** → Prometheus exporter (`/metrics`)
- **Worker** → structured logs (JSON) + Loki
- **FastAPI** → OpenTelemetry traces

9 Maintenance

Job	Schedule	Tool
nightly_dedupe	03:00	APOC <code>refactor.mergeNodes</code>
refresh_embeddings	After dedupe	Custom Python
recompute_metrics	Weekly	Neo4j Algo (betweenness)

10 Security & Backup

- **Backups** Neo4j `neo4j-admin backup` daily → `/backups/neo4j/` (btrfs snapshot).
 - **Access** Neo4j roles: **reader** for MCP; **admin** for ingestion.
 - **Secrets** Env files managed via `systemd EnvironmentFile=`.
-

11 Future Enhancements

1. Swap FAISS for **Chroma** to avoid RAM overhead.
 2. Experiment with **phi4-reasoning** as hyperedge-suggestion co-pilot.
 3. Deploy MCP in **K8s (k3s)** for HA.
-

© 2025 Bernd Prager – Licensed under Apache 2.0